

# Homework 5: ML and Kubernetes

Hrithik Dhoka

 $MS \quad CS$ 

*New York University*

## I. INTRODUCTION

The ever-growing field of machine learning (ML) relies heavily on efficient training and deployment of models. For computationally intensive tasks, distributed training approaches utilizing multiple computing resources are becoming increasingly necessary. However, managing these distributed environments can be challenging, requiring solutions that offer scalability, resource optimization, and ease of deployment.

This paper explores the use of Kubernetes, a popular container orchestration platform, for training, deploying, and performing inference on the MNIST handwritten digit classification dataset. The MNIST dataset, a well-known benchmark for image recognition, serves as a perfect example to showcase the capabilities of Kubernetes in a controlled environment.

We present a containerized workflow encompassing the entire machine learning lifecycle within a Kubernetes cluster. This includes containerizing the training script for distributed training on worker nodes, deploying the trained model as a service for real-time predictions, and utilizing Kubernetes features for efficient resource allocation and management. The code can be found [here](#)

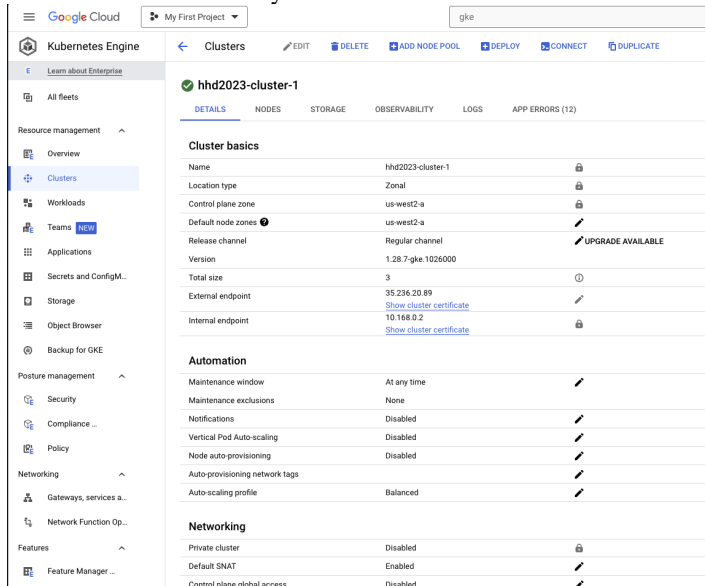
## II. METHODOLOGY

### A. Creation of cluster on GKE

We will then need to create a kubernetes cluster on Google Kubernetes Engine.

The kubectl will be configured in the cloud shell by the following command. gcloud container clusters get-credentials <kubernetes-cluster-name> --zone <zone> --project <project-name>

Now we can successfully use kubectl commands.



### B. Creating and attaching PVC resource to the cluster

- 1) To attach a PVC, we will need `pvc.yaml` which all specify all the configurations of our persistent storage.
- 2) Create a `pvc.yaml` in the cloud cluster by `vim pvc.yaml`.
- 3) In our cloud shell, we will use command `kubectl apply -f pvc.yaml` to apply these configurations to our cluster.
- 4) We can check using the command, `kubectl get pvc` or `kubectl get pvc mnist-model-pvc` whether our pvc has been successfully bounded with our cluster or not.

Thus, our PVC is ready now and successfully bound, we can go ahead with training and inference steps.

```
dh0ka@hrithiks-MBP HW5 % kubectl apply -f pvc.yaml
persistentvolumeclaim/mnist-model-pvc created
dh0ka@hrithiks-MBP HW5 % kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
mnist-model-pvc	Bound	pvc-09541962-d963-4a0e-b1cc-c46585488d9f	3Gi	RWO	standard	14s

### C. Training program in kubernetes

For training we will have to create a dockerfile, push the docker image to docker registry, and then create a training.yaml and apply that configuration to the cluster.

- 1) Build the docker image by `docker build -t mnist_training_hhd2023 . --platform=linux/amd64`. (This will name my image as `mnist_training_hhd2023`)
- 2) Then tag the image by `docker tag mnist_training_hhd2023 akohd/mnist_training_hhd2023`
- 3) Then push the train image to the dockerhub registry by using the command `docker push akohd/mnist_training_hhd2023`. Thus your train image is ready to be used by your kubernetes cluster.

- 4) Now we will create train.yaml in the cloud shell by vim train.yaml.
- 5) Deploy the training job: `kubectl apply -f train.yaml`. To check whether training is successful we can see `kubectl get jobs` and `kubectl get pods`

```

$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
inferencejob-7db99f6478-wkcdm      1/1     Running   0           23h
trainhhd2023-5nx6w                 0/1     Error     0           20m
trainhhd2023-8rvtj                 0/1     Error     0           19m
trainhhd2023-ccw4w                 0/1     Error     0           19m
trainhhd2023-fwj5k                 0/1     Error     0           14m
trainhhd2023-jpr69                 0/1     ContainerCreating   0           9m11s
trainhhd2023-kzisc                 0/1     Error     0           26m
trainhhd2023-ssck4                 0/1     Error     0           17m
trainjob-74tkc                     0/1     Completed   0           24h
trainjobhhd2023-2hp9m              0/1     Error     0           13m
trainjobhhd2023-kxm1q              0/1     Error     0           11m
trainjobhhd2023-kxhmg              0/1     Running    0           11m
trainjobhhd2023-gp866              0/1     Completed   0           10m

```

#### D. Inferencing (trained program) in kubernetes

- 1) Build the docker image by `docker build -t mnist_inference hhd2023 . --platform=linux/amd64`.

- 2) Then tag the image by docker tag mnist\_inference\_hhd2023 akohd/mnist\_inference\_hhd2023.
- 3) Then push the inference image to the dockerhub registry by using the command docker push akohd/mnist\_inference\_hhd2023.
- 4) Now, similar to training, we will create infer.yaml on our cloud shell using vim infer.yaml.
- 5) Deploy the inference application: kubectl apply -f infer.yaml
- 6) To verify whether the deployment is successful or not we can check kubectl get deployments and kubectl get pods. The status running tells that the deployment is up and can be consumed by the service.

```
dh0ka@Hrithiks-MBP HWS % kubectl apply -f infer.yaml
deployment.apps/inferjob created
dh0ka@Hrithiks-MBP HWS % kubectl get deployments
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
inferencejob  1/1     1             1           23h
inferjob      1/1     1             1           14s
```

### E. Creating Service on Kubernetes

- 1) Since our deployment is ready, we will now create a service.yaml.
- 2) Deploy the service: kubectl apply -f service.yaml.
- 3) Verify the service is running properly by kubectl get service.
- 4) Describe the service using kubectl describe service mnist-inference-service.
- 5) Now we can see our service is up and running. The external ip is the ip exposed to the end user for interacting with the application.

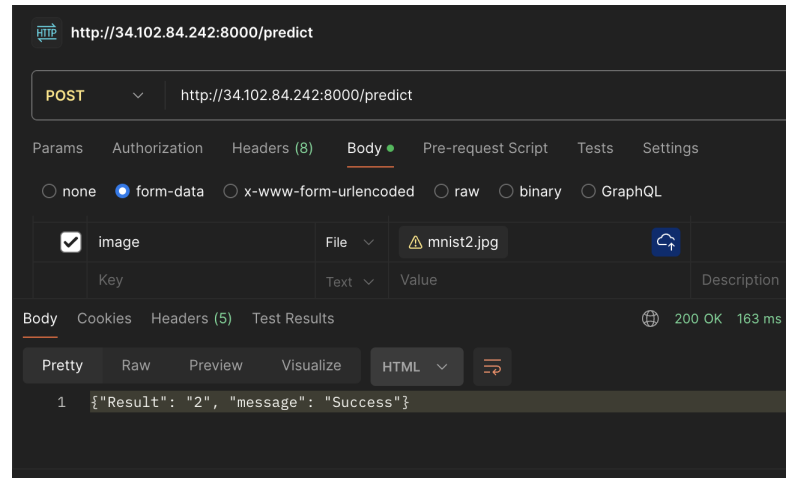
```
dh0ka@Hrithiks-MBP HWS % kubectl apply -f service.yaml
service/inferencelb created
dh0ka@Hrithiks-MBP HWS % kubectl get service
NAME          TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
inferencelb   LoadBalancer 10.24.107.39  <pending>     8000:30433/TCP   14s
kubernetes    ClusterIP      10.24.96.1    <none>         443/TCP          24h
mnist-inference-service LoadBalancer 10.24.97.75   34.102.0.130  80:32744/TCP     23h
dh0ka@Hrithiks-MBP HWS % kubectl get service
NAME          TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
inferencelb   LoadBalancer 10.24.107.39  34.102.84.242 8000:30433/TCP   3m17s
kubernetes    ClusterIP      10.24.96.1    <none>         443/TCP          24h
mnist-inference-service LoadBalancer 10.24.97.75   34.102.0.130  80:32744/TCP     23h
dh0ka@Hrithiks-MBP HWS % curl -request POST -F image=@mnist2.jpg http://34.102.84.242:8000/predict
curl: (6) Could not resolve host: POST
{"Result": "2", "message": "Success"}
```

### III. OBSERVATION

After the external ip is up and running, from commandline we use the curl command to hit the inference Flask server with an image to predict: curl -request POST -F image=@mnist2.jpg http://34.102.84.242:8000/predict. This returns a json object containing the predicted number and a message displaying success or failure depending on the input.

The API can also be tested using the Postman application to check if the external api can be accessed and returns the desired result. You can follow the steps to below for testing on postman:

- 1) Create a new POST request
- 2) Enter the external ip with the correct api route
- 3) For body, the method should accept 'form-data', with the input being a file having a key named 'image'.
- 4) Click on send, to send the request to the server and process the image
- 5) After successful execution, we get a json object containing the appropriate response based on the input.



### IV. CONCLUSION

This paper demonstrated the feasibility of using a Kubernetes cluster for training, deploying, and performing inference on the MNIST handwritten digit classification dataset. We successfully containerized the training and inference processes, leveraging the scalability and resource management capabilities of Kubernetes. The results highlight the effectiveness of this approach, showcasing distributed training for faster model development and seamless deployment for real-time predictions. Additionally, Kubernetes facilitates efficient resource utilization by dynamically allocating resources based on workload demands.

This work paves the way for exploring more complex deep learning models and real-world datasets within a Kubernetes environment. Future research directions include investigating advanced scheduling strategies for optimizing training job placement and resource allocation, as well as integrating with container orchestration tools for a more robust and automated machine learning pipeline.

### REFERENCES

- [1] Blog of deploying an ML model on Docker. <https://towardsdatascience.com/build-and-run-a-docker-container-for-your-machine-learning-model-60209c2d7a7f>
- [2] MNIST code repo=<sub>i</sub> [https://github.com/pytorch/examples/tree/main/mnist\\_rnn](https://github.com/pytorch/examples/tree/main/mnist_rnn)
- [3] Containerization wikipedia=<sub>i</sub> [https://en.wikipedia.org/wiki/Containerization\\_\(computing\)](https://en.wikipedia.org/wiki/Containerization_(computing))
- [4] Kubernetes Wikipedia =<sub>i</sub> <https://en.wikipedia.org/wiki/Kubernetes>
- [5] Kubernetes official documentation =<sub>i</sub> <https://kubernetes.io/docs/home/>
- [6] GKE official documentation =<sub>i</sub> <https://cloud.google.com/kubernetes-engine?hl=en>
- [7] Medium blog of deploying containers to kubernetes =<sub>i</sub> <https://tsai-liming.medium.com/part-3-deploying-your-data-science-containers-to-kubernetes-aac769144ec>
- [8] Gradio documentation =<sub>i</sub> <https://www.gradio.app/docs/interface>